

Mark Gameng

CS 450 – Duan Yue

Lab 2

exit() to exitStatus(int status)

Rather than just changing the existing exit system call and thereby, update all the code that used exit(), I just created a new exit call, named exitStatus(int status).

For proc.c, It's the same code as the original exit() but I just saved the exit status to the curproc. This saves the exit status of the current process.

```
void
exitStatus(int status)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    // store the exit status - Lab 2
    curproc->status = status;
```

For proc.h, I then added int status in the struct.

```
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
    int status;             // so can save exit status - Lab 2
};
```

I then declared exitStatus in user level in user.h

```
void exitStatus(int status); // lab 2
```

I also updated usys.S to have exitStatus.

```
32 | SYSCALL(exitStatus)
```

In syscall.c, I then added sys_exitStatus.

```
106 | extern int sys_exitStatus(void); // lab 2
```

```
static int (*syscalls[])(void) = {
[SYS_fork]    sys_fork,
[SYS_exit]    sys_exit,
[SYS_wait]    sys_wait,
[SYS_pipe]    sys_pipe,
[SYS_read]    sys_read,
[SYS_kill]    sys_kill,
[SYS_exec]    sys_exec,
[SYS_fstat]   sys_fstat,
[SYS_chdir]   sys_chdir,
[SYS_dup]     sys_dup,
[SYS_getpid]  sys_getpid,
[SYS_sbrk]    sys_sbrk,
[SYS_sleep]   sys_sleep,
[SYS_uptime]  sys_uptime,
[SYS_open]    sys_open,
[SYS_write]   sys_write,
[SYS_mknod]   sys_mknod,
[SYS_unlink]  sys_unlink,
[SYS_link]    sys_link,
[SYS_mkdir]   sys_mkdir,
[SYS_close]   sys_close,

// lab 2
[SYS_exitStatus] sys_exitStatus,
```

Same thing in syscall.h

```
23 | #define SYS_exitStatus 22 // lab 2
```

In sysproc.c, I just used a similar structure as sys_exit() and passed in an integer, status.

```
int
sys_exitStatus(void) // lab 2
{
    int status;

    if(argint(0, &status) < 0){
        return -1;
    }

    exitStatus(status);

    return 0; // not reached
```

I then defined it in defs.h

```
108 | void exitStatus(int); // lab 2
```

Now, exitStatus should be all implemented and what's left is testing ...

wait() -> int wait(int *status) and adding int waitpid(int pid, int *status, int options)

Now, for updating wait and adding waitpid, it's very similar to the process I did previously. Similar to exit, I made a new system call rather than updating wait, as I would have to update all the other instances of wait in xv6, which isn't that much compared to exit, but still.

For wait, I added int waitStatus, which was just the same code as the original wait, but I added code to pass back the status.

```

for(;;){
    // Scan through table looking for exited children.
    havekids = 0;
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent != curproc)
            continue;
        havekids = 1;
        if(p->state == ZOMBIE){
            // Found one.
            pid = p->pid;
            kfree(p->kstack);
            p->kstack = 0;
            freevm(p->pgdir);
            p->pid = 0;
            p->parent = 0;
            p->name[0] = 0;
            p->killed = 0;
            p->state = UNUSED;

            if(status){
                *status = p -> status; // passing back the status
            }

            release(&ptable.lock);
            return pid;
        }
    }
}

```

For waitpid, its similar to wait, but waits for a process with the given pid. Also, must wait for any process, meaning it doesn't have to be a child process.

```

for(;;){
    // Scan through table looking for exited children.
    havekids = 0;
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent != curproc)
            continue;
        if(p->pid != pid){ // if not the pid we want, skip
            continue;
        }
        if(p->state == ZOMBIE){
            // Found one.
            pid = p->pid;
            kfree(p->kstack);
            p->kstack = 0;
            freevm(p->pgdir);
            p->pid = 0;
            p->parent = 0;
            p->name[0] = 0;
            p->killed = 0;
            p->state = UNUSED;

            if(status){
                *status = p -> status; // passing back the status
            }

            release(&ptable.lock);
            return pid;
        }
    }
}

```

I also modified sysproc.c for waitStatus and waitpid.

```

int sys_waitStatus(void){
    int *status;

    if (argptr(0, (void *) &status, sizeof(*status)) < 0){
        return -1;
    }

    return waitStatus(status);
}

int sys_waitpid(void){
    int pid;
    int options;
    int *status;

    // pid is 0, status 1, options 2. cuz waitpid(pid, status, options)
    if (argint(0, &pid) < 0){
        return -1;
    }
    if (argptr(1, (void *) &status, sizeof(*status)) < 0){
        return -1;
    }
    if (argint(2, &options) < 0){
        return -1;
    }

    return waitpid(pid, status, options);
}

```

Similar to exit, for waitStatus and waitpid, I modified user.h, usys.S, syscall.c, syscall.h, defs.h

```

27 // lab 2
28 void exitStatus(int status);
29 int waitStatus(int *status);
30 int waitpid(int pid, int *status, int options);

```

```

32 SYSCALL(exitStatus)
33 SYSCALL(waitStatus)
34 SYSCALL(waitpid)

```

```

107 // lab 2
108 extern int sys_exitStatus(void);
109 extern int sys_waitStatus(void);
110 extern int sys_waitpid(void);

```

```

135 // lab 2
136 [SYS_exitStatus] sys_exitStatus,
137 [SYS_waitStatus] sys_waitStatus,
138 [SYS_waitpid] sys_waitpid,
139 };

```

```

// lab 2
#define SYS_exitStatus 22
#define SYS_waitStatus 23
#define SYS_waitpid 24

```

```

121 int wait(void);
122 int waitStatus(int*); // lab 2
123 int waitpid(int, int*, int) // lab 2

```

Testing

For testing, I first used the given usertests by xv6. Running it gave no errors and outputs “ALL TESTS PASSED”, so I then started to make my own tests for the above functions, using test.c

To be able to run the test file, I needed to update the makefile.

```

168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184     _test\

```

For test.c, this is my code.

```

printf(1, "Testing exitStatus(int status) and waitStatus(int *status)\n");
// list of childs, so i can reference their pids
//int pids[3] = {0, 0, 0};
int exitStatuses[3] = {5, 22, 13};
int pid;
int exitStatus_;

for(int i = 0; i < 3; i++){
    pid = fork();
    if (pid < 0){ // negative pid means fork failed
        printf(1, "fork failed. Retry test\n");
        exit();
    }
    //pids[i] = pid;
    if (pid == 0){ // child has their pid = 0
        printf(1, "(Child) exiting with status %d \n", exitStatuses[i]);
        exitStatus(exitStatuses[i]);
    }
}

for(int i = 0; i < 3; i++){
    pid = waitStatus(&exitStatus_);
    printf(1, "(Parent) Child with PID = %d exited with status %d \n", pid, exitStatus_);
}

```

```

printf(1, "Testing waitpid\n");
// list of childs, so i can reference their pids
int pids[3] = {0, 0, 0};
int options[3] = {7, 6, 5};

for(int i = 0; i < 3; i++){
    pid = fork();
    if (pid < 0){ // negative pid means fork failed
        printf(1, "fork failed. Retry test\n");
        exit();
    }
    if (pid == 0){ // child has their pid = 0
        printf(1, "(Child) exiting with status %d \n", exitStatuses[i]);
        exitStatus(exitStatuses[i]);
    }
    pids[i] = pid;
}

for(int i = 0; i < 3; i++){
    printf(1, "(Parent) Waiting for child with PID = %d to exit \n", pids[2 - i]);
    pid = waitpid(pids[2 - i], &exitStatus_, options[i]);
    printf(1, "(Parent) Child with PID = %d exited with status %d \n", pid, exitStatus_);
}

```

Running these tests in xv6, results in:

```
$ test
Testing exitStatus(int status) and waitStatus(int *status)
(Child) exiting with status 5
(Parent) Child with PID = 18 exited with status 5
(Child)(Child) exiting with status 13
(Parent) Child with PID = 20 exited with status 13
  exiting with status 22
(Parent) Child with PID = 19 exited with status 22
```

```
Testing waitpid
(Child) exiting with status 5
(Child(Parent) Waiting for child with PID = 23 t) exiting with status 22
(Child) exitino exit
g with status 13
(Parent) Child with PID = 23 exited with status 13
(Parent) Waiting for child with PID = 22 to exit
(Parent) Child with PID = 22 exited with status 22
(Parent) Waiting for child with PID = 21 to exit
(Parent) Child with PID = 21 exited with status 5
$ □
```

Aside from overlap from print statements, due to running at the same time, the outputs are correct and show that the methods `exitStatus(int status)`, `waitStatus(int *status)`, and `waitpid(int pid, int *status, int options)` I implemented are working.

In the picture above, the parent is waiting for child (pid = 13) to exit and doesn't run any code until it does. Once that the pid = 13 exits, can see it print above, overlapping with the parent, then the parent resumes execution and prints out that the child exited and then proceeds to wait for the next child.

I also ran `diff -r original_xv6 lab2_xv6` to show all my code changes. That file, `lab2diff.txt`, is included in the submission.