

Mark Gameng

CS 450 – Duan Yue

Lab 3 – Process Management

Change Priority

Proc.h

```
53 int prio;           // add priority value to a proc, [0, 31], lab 3
54 int runs;           // lab 3, aging? how man runs it has done. so when prio has ties, it goes by runs
55 int timestart;       // lab 3, keep track of start time
56 int burst;          // lab 3 keep track of burst, time its running
```

Proc.c

```
480 void setprio(int priority){
481     struct proc *p = myproc();
482     // prio has range of 0 and 31
483     if(priority >= 0 && priority <= 31){
484         acquire(&ptable.lock);
485         p -> prio = priority;
486         p -> state = RUNNABLE;
487         cprintf("PID: %d, Prio: %d\n", p -> pid, priority);
488         sched();
489         release(&ptable.lock);
490     }
491 }
```

Since were updating the priority of a process, we gotta lock. Also make sure that once it changes the priority value, it then transfers the control to the scheduler, sched(),. Immediately because the priority list has been updated. I also added a print statement, so its easier to see the process id and the priority value when I do the testing afterwards.

With this added system call, have to update other files

User.h

```
33 void setprio(int);
```

Defs.h

```
126 void setprio(int);
```

Usys.s

```
35 SYSCALL(setprio)
```

Syscall.c

```
113 extern int sys_setprio(void);
```

```
144 [SYS_setprio] sys_setprio,
```

Syscall.h

```
30 #define SYS_setprio 25
```

Sysproc.c

```
69 int sys_setprio(void){
70     int priority;
71     if(argint(0, &priority) < 0){
72         return -1;
73     }
74
75     setprio(priority);
76
77     return 0;
78 }
```

Initialization of priority and runs (aging) value:

Proc.c in allocproc()

```
91 p -> prio = 0; // lab 3
92 p -> runs = 0; // lab 3
93 p -> timestart = ticks; // lab 3
94 p -> burst = 0; // lab 3
```

In fork()

```
206 np -> prio = curproc -> prio; // lab 3
207 np -> runs = 0; // lab 3
208 np -> burst = 0; // lab 3
209 np -> timestart = ticks; // lab 3
```

These extra variables, runs, burst, timestart allow me to get the stats for each process so I can compare the priority and what these stats are at the end. Tracks the scheduling performance of each process as well as if my priority scheduling works or not.

Update scheduler()

```

502 struct proc *ptorun = 0; // to store the process to run // highest prio and lowest run time to avoid starvation
503 int highestprio = 32; // start off as 32 since highest is 31
504 int lowestrans = 100000; // part of aging
505
506 acquire(&ptable.lock);
507 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
508     if(p -> state != RUNNABLE)
509         continue;
510
511     if(p -> prio < highestprio && p -> runs < lowestrans){
512         highestprio = p -> prio; // keep track of highest prio
513         lowestrans = p -> runs; // keep track of lowest runs
514         ptorun = p; // store process to run later
515     }
516 }
517 // now got the highest prio and lowest run, so know which process to run first
518 // increase prio of other proc, and add a run to the proc that will run
519 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
520     if(p -> state != RUNNABLE)
521         continue;
522
523     if(p != ptorun){ // other processes not the process that we will run
524         if(p -> prio > 0){
525             p -> prio --; // increase prio, but not go to negative
526         }
527     }
528 }

```

```

532 if(ptorun){
533     // Switch to chosen process. It is the process's job
534     // to release ptable.lock and then reacquire it
535     // before jumping back to us.
536     p = ptorun; // run the process that we found to have highest prio and lowest run.
537     c->proc = p;
538     switchvm(p);
539     p->state = RUNNING;
540     p -> runs++; // add 1 to run lab 3
541     if(p -> prio < 31){
542         p -> prio++; // decrease the prio by adding 1, making sure not over 31
543     }
544     uint tmpstart = ticks; // time start
545
546     switch(&(c->scheduler), p->context);
547     switchkvm();
548
549     // Process is done running for now.
550     // It should have changed its p->state before coming back.
551     uint tmpend = ticks; // time end
552     p -> burst = p -> burst + (tmpend - tmpstart); // calc burst time and increment each time
553     c->proc = 0;
554 }
555 release(&ptable.lock);

```

Basically, the scheduler is, it goes through the list of processes and finds and stores the one with the highest priority and least amount of runs. The least amount of runs is for when there are two processes with the same priority. It will then take the one that has ran less. Once it finds the process to run, it then decreases its own priority and increments the runs var while increasing the priority of the other processes by 1. This is the aging of priority. I also calculated the total burst time, by having the ticks before starting and ticks afterwards. This is for calculating the waiting time.

In exit()

```
258 int timeend = ticks;
259 int turnaroundtime = timeend - curproc -> timestart;
260 int waitingtime = turnaroundtime - curproc -> burst;
261 cprintf("\nPID: %d, Turnaround time: %d | Waiting time: %d\n", curproc -> pid, turnaroundtime, waitingtime);
262
```

In exit, I then calculate the turnaround times and waiting time as well as printing those with the PID. The PID is so I can cross reference which process had the original priority values and see if my priority scheduling works.

Testing

For testing, I just used the code the TA showed in one of his slides.

```
11 int main(int argc, char *argv[])
12 {
13     if(argc >= 1){
14         int prio = atoi(argv[1]);
15         setprio(prio);
16         int limit = 14300;
17         int i, j;
18         for(i = 0; i < limit; i++){
19             asm("nop");
20             for(j = 0; j < limit; j++){
21                 asm("nop");
22             }
23         }
24         exit();
25     }
26 }
```

Though, I added a way for me to change the priority via the parameters when running in cmd. Also increased the limit to have a longer run time.

These were the results.

```
PID: 17, Turnaround time: 1 | Waiting time: 0
$ test 1 & ; test 10 & ; test 20

PID: 17, Turnaround time: 1 | Waiting time: 0
PID: 19, Turnaround time: 0 | Waiting time: 0
PID: 18, Prio: 1
PID: 20, Prio: 10
PID: 16, Prio: 20

PID: 18, Turnaround time: 108 | Waiting time: 71
zombie!

PID: 20, Turnaround time: 109 | Waiting time: 73
zombie!

PID: 16, Turnaround time: 115 | Waiting time: 76
```

```

PID: 39, Turnaround time: 1 | Waiting time: 0
$ test 10 & ; test 20 & ; test 1

PID: 40, Turnaround time: 1 | Waiting time: 1
PID: 41, Prio: 10

PID: 42, Turnaround time: 1 | Waiting time: 0
PID: 39, Prio: 1
PID: 43, Prio: 20

PID: 39, Turnaround time: 94 | Waiting time: 50
PID: 41, Turnaround time: 92 | Waiting time: 51
zombie!
$
PID: 43, Turnaround time: 125 | Waiting time: 86
zombie!

```

You can basically ignore the first two prints of stats, for example, PIDs 17, 19, 40, and 42. The pids with the priority is the one that matters. In both images, those with higher priority finished earlier, which is to be expected because all these processes have the same code/ same time to finish. For example, PID 39, with original priority of 1, had a turnaround time of 94 with waiting time 50, while PID 41 with originally priority of 10, had a turnaround time of 92 with a waiting time of 51. Even though, PID 39 started later than PID 41, since its priority is higher, then it ran first and ended first. The turnaround and waiting times tend to increase the lower priority it is, which is to be expected. Thus, my priority scheduling works.